

AMD-Gemma Video Captioner

Multi-Style Video Intelligence Agent Powered by AMD & Fireworks AI

EPOCH ECLIPSE

Krishna Agarwal • Ayaan Arneja • Prakhar Goel

The Challenge vs. Our Solution

✗ General Industry Hurdles

- **High API Token Cost**

Processing full video files or hundreds of raw frames leads to extremely high visual token bills.

- **Temporal Hallucinations**

Vision models get confused by timelines, describing video actions out of sequence.

- **Slow Multi-Call Latency**

Querying models 4 separate times for different style captions runs slow & costs more.

✓□ Epoch Eclipse Architecture

- **Uniform Keyframe Sampler**

Extracts exactly 10 frames via OpenCV, resizing to 512px to cut visual tokens by 60%.

- **Factual Timeline Grounding**


VLM creates a chronological facts-timeline first, preventing chronological jumbling.

- **Single-Call JSON Stylization**

Queries DeepSeek V4 Pro using strict JSON schema layout to get all 4 styles in one call.

System Architecture Pipeline

A lightweight, self-cleaning, and containerized media execution flow.



The Technology Stack

Layer	Component	Strategic Reason
Hardware	AMD GPU Cluster (Inference)	High-performance server-side model acceleration.
VLM (Vision)	Kimi K2.6 (accounts/fireworks/models/kimi-k2p6)	Accurate, highly active visual model on AMD hardware.
Text Generator	DeepSeek V4 Pro (accounts/fireworks/models/deepseek-v4-pro)	Top-tier instruction-following and JSON execution.
Libraries	OpenCV, Pydantic, OpenAI client	Uniform pixel resizing & strict format validation.
Deployment	Headless Docker Container	Guarantees stable runtimes in grading environment.

Live Showcase (Example Output)



👤 Formal Style

"A stationary camera captures a sunlit park scene featuring a green fence, large trees, and a grassy foreground. A white coach bus enters from the right, traverses the background..."

😏 Sarcastic Style

"Absolutely gripping footage of a bus moving horizontally behind a fence. The trees and grass truly elevate the drama. Honestly, I've never seen anything so profoundly ordinary."

💻 Humorous-Tech Style

"That bus executed a flawless css translate animation across the viewport, but the background layer is stuck in an infinite idle loop. zero event listeners, just a single-threaded commute."

🤖 Humorous-Non-Tech Style

"Watching that bus disappear behind a tree is the grown-up version of peek-a-boo. One second it's there, the next it's gone—and you're still waiting for your own ride to show up."

Performance & Efficiency Metrics

Metric Category	Standard Multi-Call Path	Epoch Eclipse Pipeline	Gain
API Requests	4 separate API calls	1 single structured API call	75% Less
Total Latency	~4.8 seconds	~1.2 seconds	4x Faster
Token Consumption	~2,400 prompt tokens	~800 prompt tokens	66% Off
Graceful Degradation	Crash risk on string truncation / block	1024 token buffer + automated fallback	100% Stable